Emily Xiao

emilyx@cs.cmu.edu | Website | Google Scholar | LinkedIn | Twitter

EDUCATION

M.S. Carnegie Mellon University Fall 2024 – Present
Language Technologies Institute, School of Computer Science GPA: 3.9/4.3
Advisors: Graham Neubig and Chenyan Xiong

B.A. University of California, Berkeley Fall 2017 – Spring 2021
Major: Computer Science Major GPA: 3.7/4.0

RESEARCH INTERESTS

- Large-scale machine learning (ML) and natural language processing (NLP)
- Efficient and adaptable large language models (LLMs) through methods like:
- Long-context modeling, sparse attention, data curation and synthesis, unsupervised learning

PUBLICATIONS AND PREPRINTS

Prompt-MII: Meta-Learning Instruction Induction for LLMs
Emily Xiao, *Yixiao Zeng, Ada Chen, Chin-Jou Li, Amanda Bertsch, Graham Neubig*
Preprint. Under Review.

DATE-LM: Benchmarking Data Attribution Evaluation for Large Language Models
Cathy Jiao, Yijun Pan*,* ***Emily Xiao****, Daisy Sheng, Niket Jain, Hanzhang Zhao, Ishita Dasgupta, Jiaqi W. Ma, Chenyan Xiong*
NeurIPS 2025

Efficient Many-Shot In-Context Learning with Dynamic Block-Sparse Attention
Emily Xiao, *Chin-Jou Li, Yilin Zhang, Graham Neubig, Amanda Bertsch*
ACL 2025

In-context learning with long-context models: An in-depth exploration
Amanda Bertsch, Maor Ivgi, ***Emily Xiao***, *Uri Alon, Jonathan Berant, Matthew R Gormley, Graham Neubig*
NAACL 2025 [SAC Award for Language Modeling]

Automatically generating cause-and-effect questions from passages
Katherine Stasaski, Manav Rathod, Tony Tu, ***Emily Xiao***, *Marti A Hearst*
EACL 2021, BEA Workshop

ACADEMIC RESEARCH EXPERIENCE

LLM Task Adaptation Fall 2024 – Present
CMU, Advised by Graham Neubig
- Analyzed behavior of long context in-context learning (NAACL 2025)
- Proposed efficient many-shot ICL using sparse attention (ACL 2025)
- Proposed meta-learning automatic prompt generation with RL (under submission for ICLR 2026)

LLM Training Data Curation Fall 2024 – Spring 2025
CMU, Advised by Chenyan Xiong and Jiaqi Ma
- Model-aware pre-training and post-training data selection with data attribution.
- Efficient evaluation framework and large scale analysis of existing methods (NeurIPS 2025 D&B)

LLM Systems Spring 2025
CMU, Course Project Advised by Lei Li

- Dynamic layer-wise early exit mechanism, achieving 1.4× faster inference compared to Transformer Baseline. (A+ project grade)

Synthetic Data Generation Fall 2020
UC Berkeley, Advised by Marti Hearst
- Automatically generating causal QA pairs for training and automatic evaluation (BEA@EACL 2021)

INDUSTRY EXPERIENCE

Machine Learning Engineer, TikTok Fall 2021 – Fall 2023
Query Auto-Completion
- Designed ranking model that mitigates user examination bias and click-baiting; launched globally with +2% accuracy gain. Gave talk at TikTok Search to 400+ engineers.
- Context-aware predictive pre-caching, reducing short prefix latency by >50%.
- Feature-engineered CTR ranking model.
- End-to-end optimization to capture and predict session-level query reformulation behavior.
- Integrated transformer-based query rewrite with user behavioral signals.

Related Search
- Sole developer and maintainer of the recommendation pipeline, drove 3× growth in search volume
- Boosted result diversity and freshness, designed re-ranking objectives.
- Built offline query classification model for quality and safety filtering in full ML pipeline.

Software Engineer Intern, Instagram Summer 2020
Instagram Reels – Feed Ranking (founding team)
- Developed a personalized short-video retrieval strategy; Feature-engineered CTR ranking model.

Founding Engineer, SuiteSocial Fall 2019 - Spring 2020
Startup building a brand–influencer matching platform
Won 2nd Place at UC LAUNCH Accelerator Demo Day
- Built supervised model for brand-influencer affinity ranking; Used Instragram API to extract data and features.

Founding Engineer, Prelude Spring 2019 - Summer 2019
Startup building an automated event planning platform
- Designed and built interactive webpages for MVP; conducted user interviews and product research
- Contributed to early-stage business strategy, e.g. competitor analysis and pitch preparation

TEACHING/MENTORING

Lab Assistant
Fall 2018 Data Structures & Programming Methodology (CS 61B), UC Berkeley

Student Research Mentoring

2025	Ada Chen	CMU Undergraduate
2025	Hanzhang Zhao	CMU Master

ADDITIONAL

Spoken Languages: English, Chinese, Japanese, Spanish
Coding Languages: Python, C++, SQL, Java
Tools: PyTorch, TensorFlow, NumPy, Huggingface, vLLM, SGLang, LitGPT, Hadoop, Spark, Kafka, CUDA C++
CS Coursework: Data Structures; Probability and Statistics; CS Theory; Computer Architecture; Information Devices and Systems; Computer Security; Database Systems; UI/UX Design
ML Coursework: Machine Learning; NLP; Optimization Models; Advanced NLP; Deep Learning Systems; LLM Systems; Trustworthy AI Theory and Practice; Inference Algorithms for Language Modeling